

AgentGate: Enterprise AI Assurance & Governance Fabric

Maven Poster

USE CASE

Enterprise AI agents are rapidly gaining the ability to access data, invoke tools, modify systems, and trigger business actions. As adoption scales, enterprises need a way to decide what agents are authorized to do — and to ensure the authorized action is traceable.


Tool capability is not organizational authority.

WHY IT MATTERS NOW


Sources: Gartner, 2025–26; Grand View Research, 2026.

- \$5.3B → \$24.5B enterprise agentic AI market by 2030
- <15 → 150K+ agents per average Fortune 500 enterprise by 2028.
- Only 13% of organizations say they have the right agent-governance structures in place.
- The bottleneck is shifting from agent capability to trusted governance and runtime assurance..

PROBLEM → PROTOTYPE → LEARN → ITERATE → PROVE



END-STATE ARCHITECTURE — SEPARATION OF RESPONSIBILITIES



HUMAN AUTHORITY Manager → VP Finance → CFO

EVIDENCE / PROVENANCE ActionIntent → policy → authority decision → approvals → execution receipts → verification → telemetry

Agent proposes. Authority permits. Assurance enforces, executes, and verifies. Evidence proves.

ARCHITECTURE DECISIONS

Agent paradigm	Workflow = bounded autonomous single agent	Control source	Deterministic policy outside the LLM
Runtime boundary	Typed ActionIntent + governed intercept	Human authority	VP Finance → CFO
State model	Durable workflow / checkpoint state	Assurance	Constrains + controls execution + verification
Evidence / evals	EvidenceRecord + golden state + finite-approve	Integration	Finance / external contracts + adapters

Capability does not imply authority.

Intelligence proposes. Authority permits. Assurance constrains, executes, and verifies. Evidence proves.

Model Selection Mini-Eval

Model	Avg latency	Est. cost	Result
GPT-5.4-mini	~1.96s	~\$0.0033	Baseline
GPT-5.6 Luna	~2.84s	~\$0.00086	SELECTED
GPT-5.6 Sol	larger / slower	higher	Comparison